


What We Did, Step by Step



Four steps turn a plain Apache install into a data source the collector actively monitors.

Getting Apache Ready to Be Monitored

Install Apache

sets up the httpd service on the host

Enable mod_status


exposes /server-status — a live stats page

Without mod_status, Apache has no page to read from

This part only prepares Apache itself — the collector doesn't know about it yet.

THE KEY QUESTION

Why Add Anything to agent_config.yaml?



A running service and a running collector are two separate things until the config connects them.

What Each Config Change Does

receivers: apache

defines HOW to read Apache's status page

endpoint: /server-status?auto

defines WHERE to read it from

collection_interval: 10s


defines HOW OFTEN to read it

pipelines.metrics.receivers

defines that this data joins the metrics flow

Each field answers one question: what, where, how often, and which pipeline it flows into.

Why Restart, and How We Confirmed It Worked



The config change has no effect until the collector process reloads it.