

SPLUNK DISTRIBUTION OF THE OPENTELEMETRY COLLECTOR

Built-in Receivers

A complete map of how telemetry gets into the Collector — metrics, logs, and traces

50+

Receivers

3

Pipeline types

1

Config file


Receivers get data into the Collector

Each receiver is configured, then wired into a pipeline before it becomes active.



Metrics

Host, database, infra, and cloud stats —
e.g. hostmetrics, mysql, prometheus,
vcenter



Logs

File tailing, syslog, journald, Kubernetes
events — e.g. filelog, syslog,
azure_event_hub



Traces

Distributed tracing spans — e.g. otlp,
zipkin, jaeger

Note: some receivers (otlp, kafka, splunk_hec) support multiple pipeline types at once.

Core & General-Purpose Receivers

The building blocks used in almost every Collector deployment

OTLP *otlp*

Receives OTLP data over gRPC or HTTP — metrics, logs, traces.

Host Metrics *hostmetrics*

CPU, memory, disk, network stats for the local host (agent mode).

Prometheus *prometheus*

Scrapes metrics from any Prometheus-compatible target.

Simple Prometheus *prometheus_simple*

Simplified single-target wrapper around the Prometheus receiver.

Filelog *file_log*

Tails and parses logs directly from files.

Syslog *syslog*

Parses syslog messages over TCP or UDP.

Fluent Forward *fluentforward*

TCP server accepting events via the Fluentd Forward protocol.

StatsD *statsd*

Collects StatsD messages and converts them into metrics.

TCP / UDP Log *tcp_log* / *udp_log*

Raw log ingestion directly over TCP or UDP sockets.

Journald *journald*

Parses systemd journal events (needs journalctl on PATH).

Database & Middleware Receivers

Purpose-built scrapers for common data stores and messaging systems

MySQL *mysql*

Global status and InnoDB table metrics.

PostgreSQL *postgresql*

Statistics collector queries, v9.6+.

SQL Server *sqlserver*

Microsoft SQL Server performance metrics.

Oracle DB *oracledb*

Physical reads, CPU time, and more.

MongoDB / Atlas *mongodb* / *mongodb_atlas*

Self-hosted driver stats or Atlas monitoring APIs.

Redis *redis*

Builds metrics from Redis INFO output.

Elasticsearch *elasticsearch*

Node, cluster health, and index stats.

RabbitMQ *rabbitmq*

Stats via the RabbitMQ Management Plugin.

Kafka / Kafka Metrics *kafka* / *kafkametrics*

Broker, topic, partition, and consumer group metrics.

SQL Query *sqlquery*

Runs custom SQL to generate metrics from any DB connection.

Kubernetes & Cloud Platform Receivers

For containerized and cloud-native environments

Kubelet Stats *kubeletstats*

Pod metrics from the kubelet API.

Kubernetes Cluster *k8s_cluster*

Cluster-level metrics from the Kubernetes API server.

Kubernetes Events *k8s_events*

New and updated cluster events.

Kubernetes Objects *k8sobjects*

Collects raw objects from the API server.

Receiver Creator *receiver_creator*

Instantiates other receivers at runtime via observer rules.

Discovery *discovery*

Wraps receiver creator to auto-discover collection targets.

AWS ECS Container Metrics *awsecscontainermetrics*

Task metadata and Docker stats on Amazon ECS.

Cloud Foundry *cloudfoundry*

Pulls metrics via the CF Reverse Log Proxy gateway.

Azure Event Hub *azure_event_hub*

Pulls logs from an Azure event hub.

vCenter *vcenter*

Metrics from ESXi hosts and vCenter.

Application, Network & OS Receivers

Web servers, load balancers, Windows, and infrastructure checks

Apache / Apache Spark *apache / apachespark*

Web server stats and Spark REST API metrics.

NGINX *nginx*

Stats via the ngx_http_stub_status_module endpoint.

HAProxy *haproxy*

Polls the HAProxy process via socket or HTTP.

IIS *iis*

Windows Performance Counters for IIS.

Windows Event Log / Perf Counters *windows_event_log / windowsperfcounters*

Windows-only log and counter collection.

Active Directory DS *active_directory_ds*

Domain controller metrics via Performance Counters.

HTTP Check / SSH Check *http_check / sshcheck*

Synthetic checks against HTTP or SSH endpoints.

Chrony *chrony*

Time-sync stats — a Go reimplement of chronyc.

Zookeeper *zookeeper*

Metrics from a Zookeeper instance.

JMX *jmx*

MBean server metrics via the OTel JMX Metric Gatherer.

Splunk-Native & Specialty Receivers

Receivers unique to the Splunk distribution, plus niche protocols

Splunk HEC *splunkhec*

Accepts telemetry in Splunk HTTP Event Collector format.

Splunk Enterprise *splunkenterprise*

Ingests operational health metrics from a Splunk deployment.

Smart Agent *smartagent*

Runs legacy SignalFx Smart Agent monitors inside the Collector.

SignalFx Gateway Prom. Remote Write

signalfxgatewayprometheusremotewritereceiver

Native OTel version of the SFX remote-write gateway.

Zipkin / Jaeger *zipkin / jaeger*

Trace ingestion for Zipkin v1/v2 and Jaeger formats.

Carbon / Wavefront *carbon / wavefront*

Carbon plaintext protocol; Wavefront reuses its transport.

Snowflake *snowflake*

Account metrics via Snowflake queries.

Pure Storage FlashArray *purefa*

Metrics from Pure Storage FlashArray.

Cisco OS *cisco_os*

Cisco Nexus switch metrics over SSH.

CollectD *collectd*

Data from the CollectD write_http plugin (JSON only).

A Naming Convention Is In Flight

OpenTelemetry-upstream is renaming multi-word receivers to snake_case. The Splunk distribution already conforms to these new names.

azureblob → **azure_blob**

azureeventhub → **azure_event_hub**

azuremonitor → **azure_monitor**

filelog → **file_log**

httpcheck → **http_check**

mongodbatlas → **mongodb_atlas**

prometheusremotewrite → **prometheus_remote_write**

tcplog → **tcp_log**

tlscheck → **tls_check**

udplog → **udp_log**

windowseventlog → **windows_event_log**

yanggrpc → **yang_grpc**

Key Takeaways

1

Nothing runs until it's piped

Receivers are declared under `receivers;`, but only take effect once referenced in a `service.pipelines` block.

2

Start with OTLP + Host Metrics

Nearly every agent-mode config pairs `otlp` (app telemetry) with `hostmetrics` (infra telemetry).

3

Reach for Smart Agent or Discovery for legacy or dynamic targets

`smartagent` covers hundreds of legacy monitors; `receiver_creator` / `discovery` auto-detect endpoints in containers and Kubernetes.

4

The Splunk list can lag upstream

For the newest or niche components, check `opentelemetry-collector-contrib` directly.